

Requisite packages and data:

```
library(dplyr)

# data
library(AmesHousing)

# Load the housing data, clean names, then grab just six columns
# to simplify the dataset for display purposes.
ames_data <- make_ames() %>%
  janitor::clean_names() %>%
  select(sale_price, bsmt_fin_sf_1, first_flr_sf,
         total_bsmt_sf, neighborhood, gr_liv_area)
```



Superceding Functions

There are two major families of functions that supercede old functionality

The Replacements:

- ``across()``

- ``slice()``

In addition, there are some deprecated functions that are worth noting as well as some new `mutate()` arguments. In this post, I'll walk through some examples of each of these changes.

Across

All `*_if()`, `*_at()` and `*_all()` function variants were superseded in favor of `across()`. `across()` makes manipulating multiple columns more intuitive and consistent with other dplyr syntax.

`across()` is my favorite new dplyr function because I've always had to stop and think and pull up the docs when using `mutate_if()` and `mutate_at()`. Most notably, I appreciate the use of tidy selection rather than the `vars()` method used in `mutate_at()`.

Let's see `across()` in action. Let's say we want to convert all the square foot variables to square yards. When we take a look at our data, we see that all of the square foot variables either contain "area" or "_sf" in their names.

```
feet_to_yards <- function(x) {x / 9}
```

Here is the old way this was done with `mutate_at()`:

```
ames_data %>%
  mutate_at(.vars = vars(contains("_sf") | contains("area")), .funs =
    feet_to_yards)
```

Now we use `across()` in combination with a vector. In this case, we used `contains()` to grab variable names that contain "_sf" or "area".

```
ames_data %>%
  mutate(across(.cols = c(contains("_sf"), contains("area")), .fns =
    feet_to_yards)) %>%
  head()
```

```
## # A tibble: 6 x 6
##   sale_price bsmt_fin_sf_1 first_flr_sf total_bsmt_sf neighborhood
##   gr_liv_area
##
## 1      215000          0.222         184          120 North_Ames      184
## 2      105000          0.667          99.6           98 North_Ames
## 3      172000          0.111         148          148. North_Ames      148.
## 4      244000          0.111         234          234. North_Ames      234.
## 5      189900          0.333         103          103. Gilbert        181
## 6      195500          0.333         103          103. Gilbert        178.
```

`across()` can also replace `mutate_if()` in combination with `where()`.

Old way with `mutate_if()`:

```
ames_data %>%
  mutate_if(is.numeric, log)
```

New way with `across(where())`:

```
## new dplyr(log transform numeric values)
ames_data %>%
  mutate(across(where(is.numeric), log)) %>%
  head()
```

```
## # A tibble: 6 x 6
```

```
##   sale_price bsmt_fin_sf_1 first_flr_sf total_bsmt_sf neighborhood
gr_liv_area
##
## 1      12.3      0.693      7.41      6.98 North_Ames
7.41
## 2      11.6      1.79      6.80      6.78 North_Ames
6.80
## 3      12.1      0      7.19      7.19 North_Ames
7.19
## 4      12.4      0      7.65      7.65 North_Ames
7.65
## 5      12.2      1.10      6.83      6.83 Gilbert
7.40
## 6      12.2      1.10      6.83      6.83 Gilbert
7.38
```

`summarize()` now uses the same `across()` and `where()` syntax that we used above with `mutate`. Let's find the average of all numeric columns for each neighborhood.

```
ames_data %>%
  group_by(neighborhood) %>%
  summarize(across(where(is.numeric), mean, .names = "mean_{col}")) %>%
  head()

## `summarise()` ungrouping output (override with `.groups` argument)

## # A tibble: 6 x 6
##   neighborhood mean_sale_price mean_bsmt_fin_s... mean_first_flr_...
##
## 1 North_Ames      145097.      3.66      1175.
## 2 College_Cre...    201803.      4.01      1173.
## 3 Old_Town        123992.      5.80      945.
## 4 Edwards         130843.      4.27      1115.
## 5 Somerset        229707.      4.59      1188.
## 6 Northridge_...    322018.      3.99      1613.
## # ... with 2 more variables: mean_total_bsmt_sf , mean_gr_liv_area
```

As you can see, we calculated the neighborhood average for all numeric values. On top of that, we were able to easily prefix the column names with “mean_” thanks to another useful `across()` argument called “.names”.

Not only that, in conjunction with the `where()` helper, `across()` unifies “_if” and “_at” semantics, allowing more intuitive and elegant column selection. For example, let's mutate the square footage variables that are integers (like `mutate_if()`), and the square footage variables that end with “_sf” (like `mutate_at()`) to make them doubles.

```
ames_data %>%
  mutate(across(where(is.integer) & ends_with("_sf"), as.double))

## # A tibble: 2,930 x 6
##   sale_price bsmt_fin_sf_1 first_flr_sf total_bsmt_sf neighborhood
gr_liv_area
##
## 1      215000      2      1656      1080 North_Ames
1656
## 2      105000      6      896      882 North_Ames
896
## 3      172000      1      1329      1329 North_Ames
1329
## 4      244000      1      2110      2110 North_Ames
```

```

2110
## 5      189900      3      928      928 Gilbert
1629
## 6      195500      3      926      926 Gilbert
1604
## 7      213500      3     1338     1338 Stone_Brook
1338
## 8      191500      1     1280     1280 Stone_Brook
1280
## 9      236500      3     1616     1595 Stone_Brook
1616
## 10     189000      7     1028      994 Gilbert
1804
## # ... with 2,920 more rows

```

Notice, the “first_flr_sf” was converted to a double, but the “gr_living_area” remains an integer because it doesn’t fit the criteria `aends_with(“_sf”)`.

`across()` can also perform `mutate_all()` functionality with `across(everything(), ...`

Slice

`top_n()`, `sample_n()`, and `sample_frac()` have been superseded in favor of a new family of `slice()` helpers.

Reasons for future deprecation:

- `top_n()` – has a confusing name that might reasonably be considered to filter for the min or the max rows. For example, let’s say we have data for a track and field race that records lap times. One might reasonable assume that `top_n()` would return the fastest times but they actually return the times that took the longest. `top_n()` has been superseded by `slice_min()`, and `slice_max()`.
- `sample_n()` and `sample_frac()` – it’s easier to remember (and pull up documentation for) two mutually exclusive arguments to one function called `slice_sample()`.

```

# Old way to grab the five most expensive homes by sale price
ames_data %>%
  top_n(n = 5, wt = sale_price)

# New way to grab the five most expensive homes by sale price
ames_data %>%
  slice_max(sale_price, n = 5)

## # A tibble: 5 x 6
##   sale_price bsmt_fin_sf_1 first_flr_sf total_bsmt_sf neighborhood
gr_liv_area
##
## 1      755000      3      2444      2444 Northridge
4316
## 2      745000      3      2411      2396 Northridge
4476
## 3      625000      3      1831      1930 Northridge
3627
## 4      615000      3      2470      2535 Northridge_He...
2470
## 5      611657      3      2364      2330 Northridge_He...
2364

# You can also grab the five cheapest homes
ames_data %>%

```

```

    slice_min(sale_price, n = 5)

## # A tibble: 5 x 6
##   sale_price bsmt_fin_sf_1 first_flr_sf total_bsmt_sf neighborhood
gr_liv_area
##
## 1      12789           7          832          678 Old_Town
832
## 2      13100           5          733           0 Iowa_DOT_and_...
733
## 3      34900           6          720          720 Iowa_DOT_and_...
720
## 4      35000           7          498          498 Edwards
498
## 5      35311           2          480          480 Iowa_DOT_and_...
480

# Old way to sample four random rows(in this case properties)
ames_data %>%
  sample_n(4)

# New way to sample four random rows(in this case properties)
ames_data %>%
  slice_sample(n = 4)

## # A tibble: 4 x 6
##   sale_price bsmt_fin_sf_1 first_flr_sf total_bsmt_sf neighborhood
gr_liv_area
##
## 1      119000           6          948          948 Edwards
948
## 2      156000           1          990          990 College_Creek
990
## 3      245700           3         1614         1614 Northridge_He...
1614
## 4      108000           2         1032          1032 Old_Town
1032

# Old way to sample a random 0.2% of the rows
ames_data %>%
  sample_frac(0.002)

# New way to sample a random 0.2% of the rows
ames_data %>%
  slice_sample(prop = 0.002)

## # A tibble: 5 x 6
##   sale_price bsmt_fin_sf_1 first_flr_sf total_bsmt_sf neighborhood
gr_liv_area
##
## 1      110000           7          682          440 Old_Town
1230
## 2      136000           6         1040         1040 North_Ames
1040
## 3      208000           1         1182          572 Crawford
1982
## 4      115000           1          789          789 Old_Town
789
## 5      145500           1         1053         1053 North_Ames

```

Additionally, `slice_head()` and `slice_tail()` can be used to grab the first or last rows, respectively.

Nest By

`nest_by()` works similar to `group_by()` but is more visual because it changes the structure of the tibble instead of just adding grouped metadata. With `nest_by()`, the tibble transforms into a rowwise dataframe (Run `vignette("rowwise")` to learn more about the revised rowwise functionality in dplyr 1.0.0).

First, for the sake of comparison, let's calculate the average sale price by neighborhood using `group_by()` and `summarize()`:

```
ames_data %>%
  group_by(neighborhood) %>%
  summarise(avg_sale_price = mean(sale_price)) %>%
  ungroup() %>%
  head()

## `summarise()` ungrouping output (override with `.groups` argument)

## # A tibble: 6 x 2
##   neighborhood      avg_sale_price
##
## 1 North_Ames          145097.
## 2 College_Creek       201803.
## 3 Old_Town            123992.
## 4 Edwards            130843.
## 5 Somerset            229707.
## 6 Northridge_Heights    322018.
```

The `summarize()` operation works well with `group_by()`, particularly if the output of the summarization function are single numeric values. But what if we want to perform a more complicated operation on the grouped rows? Like, for example, a linear model. For that, we can use `nest_by()` which stores grouped data not as metadata but as lists in a new column called "data".

```
nested_ames <- ames_data %>%
  nest_by(neighborhood)

head(nested_ames)

## # A tibble: 6 x 2
## # Rowwise: neighborhood
##   neighborhood      data
##
## 1 North_Ames          >
## 2 College_Creek      [443 x 5]
## 3 Old_Town            [267 x 5]
## 4 Edwards            [239 x 5]
## 5 Somerset            [194 x 5]
## 6 Northridge_Heights [182 x 5]
## 6 Northridge_Heights [166 x 5]
```

As you can see, `nest_by()` fundamentally changes the structure of the dataframe unlike `group_by()`. This feature becomes useful when you want to apply a model to each row of the nested data.

For example, here is a linear model that uses square footage to predict sale price applied to each neighborhood.

```
nested_ames_with_model <- nested_ames %>%
  mutate(linear_model = list(lm(sale_price ~ gr_liv_area, data = data)))
```

```
head(nested_ames_with_model)

## # A tibble: 6 x 3
## # Rowwise: neighborhood
##   neighborhood                data linear_model
##           >
## 1 North_Ames                [443 x 5]
## 2 College_Creek            [267 x 5]
## 3 Old_Town                  [239 x 5]
## 4 Edwards                   [194 x 5]
## 5 Somerset                  [182 x 5]
## 6 Northridge_Heights        [166 x 5]
```

It's important to note that the model must be vectorized, a tranformation performed here with list(). Let's take a look at the model that was created for the "North_Ames" neighborhood.

```
north_ames_model <- nested_ames_with_model %>%
  filter(neighborhood == "North_Ames") %>%
  pull(linear_model)

north_ames_model

## [[1]]
##
## Call:
## lm(formula = sale_price ~ gr_liv_area, data = data)
##
## Coefficients:
## (Intercept)  gr_liv_area
##      74537.97      54.61
```

The model shows that for each additional square foot, a house in the North Ames neighborhood is expected to sell for about \$54.61 more.

Additional Mutate arguments

Control what columns are retained with ".keep"

```
# For example "used" retains only the columns involved in the mutate
ames_data %>%
  mutate(sale_price_euro = sale_price / 1.1, .keep = "used") %>%
  head()

## # A tibble: 6 x 2
##   sale_price sale_price_euro
##
## 1      215000      195455.
## 2      105000       95455.
## 3      172000     156364.
## 4      244000     221818.
## 5      189900     172636.
## 6      195500     177727.
```

Control where the new columns are added with ".before" and ".after"

```
# For example, make the "sale_price_euro" column appear to the left of the
"sale_price" column like this
ames_data %>%
```

```
mutate(
  sale_price_euro = sale_price / 1.1, .keep = "used", .before = sale_price
) %>%
head()

## # A tibble: 6 x 2
##   sale_price_euro sale_price
##
## 1          195455.      215000
## 2           95455.      105000
## 3        156364.      172000
## 4        221818.      244000
## 5        172636.      189900
## 6        177727.      195500
```

Conculsion

This was a short, high level look at my favorite new features coming in dplyr 1.0.0. The two major changes were the addition of `across()` and `slice()` which supercede old functionality. `across()` makes it easy to mutate specific columns or rows in a more intuitive, consistent way. `slice()` makes similar improvements to data sampling methods. I am also a big fan of the new `nest_by()` functionality, and plan to search for elegant ways to incorporate it in my upcoming R projects. These changes align dplyr syntax more closely with conventions common in the tidyverse. Thanks tidyverse team for continually pushing the boundaries to make data analytics easier in practice and to learn/teach!